

PhysioFlow: an open-source interactive platform for crop ecophysiology and agronomic data analysis with statistical and experimental-design guard-rails

Daniel Hilário da Silva^{a,*}, Leandro Rodrigues da Silva Souza^{b,c}, André Abade^d, Francisco de Almeida Lobo^e, Fernando Rodrigues Cabral Filho^{b,c}, Erli Pinto dos Santos^f, Adinan Alves da Silva^c, Daiane Alves da Silva^b, Ketlyn Santos Sousa^c, Igor Eli da Silva^c, Alan Carlos da Costa^c

^aInstituto Federal Goiano, Campus Catalão, Goiás, Brazil

^bCentro de Excelência em Agricultura Exponencial (CEAGRE), Rio Verde, Goiás, Brazil

^cInstituto Federal Goiano, Campus Rio Verde, Goiás, Brazil

^dInstituto Federal de Mato Grosso, Cuiabá, Mato Grosso, Brazil

^eUniversidade Federal de Mato Grosso, Mato Grosso, Brazil

^fUniversidade Federal Rural do Rio de Janeiro, Seropédica, Rio de Janeiro, Brazil

*Corresponding author: daniel.hilario@ifgoiano.edu.br

ORCID: Daniel Hilário da Silva: 0000-0002-0800-065X; Leandro Rodrigues da Silva Souza: 0000-0003-2477-6893; André Abade: 0000-0001-9771-9123; Francisco de Almeida Lobo: 0000-0002-5670-0351; Fernando Rodrigues Cabral Filho: 0000-0002-5090-5946; Erli Pinto dos Santos: 0000-0002-4888-906X; Adinan Alves da Silva: 0000-0002-1589-6773; Daiane Alves da Silva: 0009-0004-8550-6576; Ketlyn Santos Sousa: 0009-0001-5682-3925; Igor Eli da Silva: 0000-0001-8897-152X; Alan Carlos da Costa: 0000-0001-9339-9257.

Abstract

PhysioFlow is an open-source Python/Streamlit platform for crop-ecophysiology and general agronomic or biological data. It ingests CSV/TSV/XLSX exports, validates them against an explicit schema, and offers a configurable replicate pipeline, exploratory analysis with statistical guard-rails (confounding, multicollinearity, normality and multi-method outlier checks), supervised regression and classification with optional GroupK-Fold, geospatial analysis in UTM metres, STL decomposition, and a formal experimental-design module (from completely randomised and block designs to factorial, Latin square, split-plot, strip-plot, nested, ANCOVA and dose-response), with selected procedures checked against reference analyses in R. An automatic data-profile makes it dataset-agnostic beyond ecophysiology. Released under GPL-2.0-or-later, it runs code-free in a trilingual (PT/EN/ES) interface.

Keywords: Crop ecophysiology; IRGA; Exploratory data analysis; Experimental design; Statistical guard-rails; Reproducible research

Code metadata

Nr.	Code metadata description	Value
C1	Current code version	v1.0 (commit bc69e10)
C2	Permanent link to code/repository used for this code version	https://github.com/ML-Carbon-Project/physioFlow/tree/bc69e105774048526480fe9bfab633fbd5527e3f
C3	Permanent link to Reproducible Capsule	https://doi.org/10.24433/CO.9901925.v1
C4	Legal code license	GNU General Public License v2.0 or later (GPL-2.0-or-later)
C5	Code versioning system used	git
C6	Software code languages, tools and services used	Python 3.10+; Streamlit; pandas, NumPy, SciPy, scikit-learn, statsmodels; geopandas, shapely, pyproj, libpysal, esda; pytest
C7	Compilation requirements, operating environments & dependencies	Linux, macOS, Windows; Python ≥ 3.10 ; <code>pip install -r requirements.txt</code>
C8	If available, link to developer documentation/manual	https://github.com/ML-Carbon-Project/physioFlow/tree/main/docs
C9	Support email for questions	daniel.hilario@ifgoiano.edu.br

Table 1: Code metadata.

1 Introduction

Plant ecophysiology links the carbon, water and energy economy of crops to their environment, and its field measurement underpins decisions in agronomy, breeding and climate-change adaptation [1, 2]. Three portable instruments dominate field campaigns: infrared gas analyzers (IRGA) measure net photosynthesis (A), transpiration (E), stomatal conductance (g_s) and internal/ambient CO_2 , from which derived efficiencies such as WUE, A/C_i and C_i/C_a are computed [3–5]; chlorophyll fluorometers and content meters report Y_{II} , the electron transport rate and chlorophyll a/b [6, 7]; and ceptometers estimate the leaf area index [8]. Together they characterise crop photosynthetic performance across genotypes, phenological stages and seasons, anchoring studies of responses to irrigation, salinity and soil management [9–13], central in the Brazilian Cerrado to both yield and ecosystem carbon accounting [14].

Each instrument exports a wide CSV/XLSX table whose rows are leaf or canopy measurements with gas-exchange, fluorescence, chlorophyll and LAI variables plus metadata (crop, farm, land use, season, coordinates, date). Turning this raw output into a publishable analysis still demands several non-trivial steps: standardising categorical fields, discarding rows without minimum metadata, consolidating chlorophyll and LAI replicates, running exploratory and inferential analyses, checking multicollinearity among derived indices, inspecting spatial structure, decomposing temporal trends and, optionally, fitting predictive models.

Despite the wide availability of these instruments, manufacturers provide no free, open-source software for the exploratory and modelling stages that turn raw exports into results; vendor viewers focus on acquisition and quality control. Researchers must therefore write custom scripts in R [15] or Python, a considerable barrier for those without programming or statistics training. In most laboratories these steps are still done by hand or half-automated across spreadsheets and disconnected packages, which is slow, error-prone, hard to reproduce, and, in high-throughput campaigns, the bottleneck of the whole analysis. The structure of ecophysiology field data also invites silent statistical errors: physiological indices are mutually derived (e.g. $\text{WUE} = A/E$), inflating multicollinearity by construction; sampling designs frequently alias farm, crop and land use, so farm and crop effects cannot be estimated separately; replicated measurements at the same site induce pseudoreplication [16] that inflates model skill; and short campaigns offer too few dates for meaningful seasonal decomposition.

Inspired by recent open-source efforts that bridged data analysis and end users via Streamlit [17–19], we developed *PhysioFlow* to consolidate the entire ecophysiology data-analysis workflow into a single configurable,

interactive platform. The distinguishing design choice is a set of *statistical guard-rails* woven into the user interface: the application not only computes statistics but warns when their assumptions are violated. Recent software papers in neighbouring agronomic niches [20, 21] have shown how much is gained by releasing an instrument-bound, end-to-end analytical pipeline as citable software, and *PhysioFlow* does the same for crop ecophysiology.

2 Software description

2.1 Software architecture

PhysioFlow is implemented in Python 3.10 or later [17, 22, 23] and follows a session-state, page-renderer pattern in which each of nine user-facing pages reads from a shared `st.session_state` and consumes the cleaning, statistics and visualisation primitives in `src/`, so costly I/O and cleaning happen at most once per session (Fig. 1). A resolver (`src/profile.py`) matches the uploaded frame against the reference schema and selects the *data-profile*: the physiology profile (domain defaults, presets and schema report) when the columns match, otherwise a generic profile that drops those presets so any agronomic or biological table is analysed without code changes.




Figure 1: High-level data-flow architecture of *PhysioFlow*. Raw exports from IRGA, chlorophyll fluorometer, and ceptometer instruments are validated against the reference schema upon upload, cached in the Streamlit session state, and consumed by the Pipeline, which produces a single cleaned frame. A shared, reactive *global filter panel* (crop, farm, municipality, season, date range, and replicate mode) is injected at the top of the four analytical pages (EDA, Regression, Modelling, Spatial) and propagates instantly to every chart, test and model; the Time Series, Group-comparison and Experimental-design pages read the cleaned frame directly through their own controls. Library stack: Pipeline → `pandas/NumPy`; modelling → `scikit-learn` [24]; inference → `SciPy/statsmodels`; visualisation → `seaborn` [25, 26]; geospatial → `geopandas, pyproj, libpysal/esda` [27].

The `src/` package is organised by concern: `schema.py` validates 31 reference columns in three severity tiers without blocking the upload; `pipeline.py` holds the deterministic cleaning primitives; `stats_utils.py` implements bias-corrected Cramér’s V [28], the confounding auditor and the analysis-of-variance routines behind the experimental-design page (split-plot, strip-plot, nested and ANCOVA models; Scott–Knott, Duncan, LSD, Scheffé and Dunnett comparisons), each checked against reference implementations in R; and `ml/model_registry.py` declares two `scikit-learn` pipelines of five regressors [29, 30] and seven classifiers. Cross-cutting `i18n/` (775 keys at full PT/EN/ES parity), `auth/`, `config/` and `tests/` (113 `pytest` functions, 169 test cases) serve every page. Public functions carry type annotations, and a GitHub Actions workflow

runs `ruff`, `mypy` and the `pytest` suite on Python 3.10, 3.11 and 3.12 on every push to the main branch; new estimators, cleaning primitives and schema columns are registered declaratively.

2.2 Software functionalities

Upload (Fig. 2). Accepts CSV, TSV, TXT (configurable delimiter), XLSX and XLS files. The schema validator classifies every expected column as *required*, *recommended* or *optional*, checks types, flags columns that exist but are 100% empty, and verifies that latitude/longitude lie within plausible ranges. Missing columns disable dependent features without blocking the upload.




Figure 2: Upload page after ingestion of an ecophysiology dataset. The schema validator reports each expected column with its severity tier, the resolved column name, the detected type and the dependent feature.

Pipeline. Four deterministic steps (text standardisation, removal of rows without essential metadata, removal of empty grid points and replicate treatment) are followed by a step-by-step report (Fig. 3) recording rows in/out

per step, with a warning whenever a step discards more than 50% of the rows. Chlorophyll and LAI replicates are handled in six modes: mean, median, replicate expansion into separate rows (which does not, by itself, make them statistically independent observations), and exclusive selection of replicate 1, 2 or 3.

Pipeline & Processing

Automated Physiology Cleaning Pipeline

This pipeline runs the following structured steps:

- Standardization:** Trimming of extra whitespace in text columns.
- Metadata filter:** Removal of rows lacking essential fields (Cultura, Uso atual, Época).
- Empty-grid filter:** Removal of sampling-grid rows with no actual measurements.
- Replicate handling:** Consolidation or unfolding of Chlorophyll and LAI according to the selected mode.

Change replicate treatment:

Mean of Replicates (Grouped) ▾

Original rows

Processed rows

1661

81

The pipeline removed 95.1% of rows (1661 → 81). Check whether essential columns (Cultura, Uso atual, Época) are filled in the spreadsheet.

Step report

Step "Removal of records without essential metadata" removed 94.9% of rows (1661 → 85). May indicate large-scale missing data — investigate before proceeding.

	Step	Rows before	Rows after	Removed	% removed
0	Text standardization (string stripping)	1661	1661	0	0
1	Removal of records without essential metadata	1661	85	1576	94.8826
2	Removal of empty grid points (no measurements)	85	81	4	4.7059
3	Replicate consolidation by arithmetic mean	81	81	0	0

Figure 3: Step-by-step audit of the cleaning pipeline applied to the de-identified Rio Verde dataset. Each row reports lines before/after and the share removed; the essential-metadata filter alone removes 94.9% of the raw rows, triggering the high-discard warning.

Exploratory Data Analysis. Twelve tabs cover descriptive statistics; a data-quality tab auditing categorical *confounding* via Cramér’s V (Fig. 4); distributions, boxplots, pair plots and Pearson/Spearman/Kendall correlations; a lat/lon map; temporal aggregation and composition charts; non-parametric inference (Kruskal–Wallis [31] with a minimum-N slider, Shapiro–Wilk [32] with Q–Q plots, and VIF [33] caveating that derived indices such as C_i/C_a , A/C_i and WUE inflate VIF by construction); a top- N hotspots ranking; and a multi-method outlier audit (Z-score, IQR, Isolation Forest [34], Local Outlier Factor [35], Elliptic Envelope) with a consensus criterion.

Regression and Modelling. The Regression page offers physiological bivariate presets (e.g. g_s vs. A) and free custom regression. On the Modelling page the user chooses the task: regression compares five estimators and reports holdout R^2 , MAE, RMSE and cross-validated R^2 ; classification compares seven and reports accuracy,

Confounding between categories

Pairs of categorical columns that partition rows the same way. When two columns are redundant (e.g. *Fazenda* $\leftrightarrow$ *Cultura*), any effect attributed to one is statistically indistinguishable from the other — read Moran's I, Kruskal-Wallis and group comparisons cautiously.

	Column A	Column B	Cramér's V	Relation	n
0	Fazenda	Uso atual	1	Redundant ($A \equiv B$)	81
1	Fazenda	Cultura	1	Redundant ($A \equiv B$)	81
2	Uso atual	Cultura	1	Redundant ($A \equiv B$)	81
3	Fazenda	Data da coleta	0.994	Data da coleta determines Fazenda	81
4	Uso atual	Data da coleta	0.994	Data da coleta determines Uso atual	81
5	Cultura	Data da coleta	0.994	Data da coleta determines Cultura	81
6	Época	Data da coleta	0.994	Data da coleta determines Época	81
7	Fazenda	Estágio	0.941	Estágio determines Fazenda	81
8	Uso atual	Estágio	0.941	Estágio determines Uso atual	81
9	Cultura	Estágio	0.941	Estágio determines Cultura	81

There are fully redundant pairs (one column is just a relabeling of the other). Do not use both as independent factors in models or comparisons.

Figure 4: Data-quality tab auditing categorical confounding. On the Rio Verde dataset the tool detects that *Fazenda* (farm), *Cultura* (crop) and *Uso atual* (land use) are perfectly redundant (Cramér's $V = 1.00$), warning the user that farm, crop and land-use effects cannot be identified separately from this sampling design.

precision, recall, F_1 and a confusion matrix. In both, an *optional GroupKFold* with a grouping-column selector (by default a synthetic *farm + point* key) curbs pseudoreplication [16].

Experimental Design. A separate page treats the uploaded table as a designed experiment: the user names a categorical treatment factor and a numeric response and chooses the layout, from completely randomised (CRD) and randomised complete block (RCBD) through two- and three-factor factorials to the Latin square and the compound-error split-plot, strip-plot and nested designs; analysis of covariance (ANCOVA) and dose-response regression are also available. The page fits the model, prints the ANOVA table, checks residual normality and homoscedasticity (Shapiro–Wilk, Levene) on screen, and warns on violated assumptions or under-replicated groups. Mean comparison offers Tukey HSD, Scott–Knott [36], Duncan, LSD/Fisher, Scheffé and Dunnett as compact-letter displays, with a reproducible code snippet per result. The design routines were checked against reference analyses in R on the same datasets and model specifications: split-plot, strip-plot and nested ANOVA with their design-specific error strata, and Type II tests from `car::Anova` for the factorial and ANCOVA examples, with adjusted means compared against `emmeans`. In the published capsule (C3), 31 of these quantities are recomputed against R 4.4.2 running inside the reproducible run itself, the largest absolute difference being 6.7×10^{-11} . The mean-comparison procedures were checked against `stats::TukeyHSD`, the CRAN `ScottKnott` package, `agricolae` (LSD, Duncan and Scheffé) and `multcomp` (Dunnett), reproducing the same pairwise decisions on public datasets; the comparison against `multcomp` was restricted to one-way examples. Reference values are frozen as regression tests in `tests/test_sample_datasets.py`, so the suite guards them without requiring R at continuous-integration time, and the comparisons are tabulated in the repository (`docs/validacao_externa.md`). The Latin square and dose-response routines are covered by internal tests only.

Spatial Analysis. IDW interpolation, global Moran's I and LISA cluster maps [37], Getis–Ord G_i^* [38], UTM grid aggregation and ordinary kriging [39] with a fitted spherical variogram. All distance-based computations run in *UTM metres*: the zone is selected from the mean longitude and the hemisphere from the mean latitude,

so distances are computed in metres rather than from degrees of latitude and longitude.

Time Series. Daily resampling (mean or median) and STL decomposition [40], *blocked when fewer than ten distinct dates are present* so that sparse campaigns do not yield a decomposition dominated by interpolation.

Two-group comparison and Authentication. A comparative page splits the dataset by multi-select or substring pattern and returns a per-variable Mann–Whitney U test [41], paired log-linear regressions and an hourly profile. An optional Supabase login (off by default) adds authentication.

3 Results and discussion

3.1 A crop-ecophysiology campaign

A de-identified Cerrado campaign from Rio Verde, GO, with 1 661 raw measurement rows across soybean and sugarcane, was ingested via the Upload page. The schema validator confirmed the required metadata and physiological variables and flagged *Manejo* (soil management) and *Textura* (soil texture) as present but 100% empty.

The Pipeline was run in the default mean-replicate mode. Its audit trail (Fig. 3) traces 1 661 input rows down to 81 sampling-point records: the essential-metadata filter alone removes 94.9% of the raw rows (an incomplete sampling grid in which most planned points carry no measurement), triggering the high-discard warning before any inference is attempted.

The cleaned frame was then submitted to the EDA data-quality tab (Fig. 4). The Cramér’s V auditor reported that *Fazenda* (farm), *Cultura* (crop) and *Uso atual* (land use) are mutually redundant (Cramér’s V = 1.00): the two surviving farms map one-to-one onto the two crops (Farm A $\equiv$ sugarcane, 27 records; Farm B $\equiv$ soybean, 54 records). The guard-rail makes explicit that farm and crop effects cannot be estimated separately from these data, because each crop occurs on only one farm; a plain grouped comparison would quietly read the difference as spatial structure instead.

The Modelling page then predicted net photosynthesis (A , $n = 71$) from gas-exchange, fluorescence and chlorophyll predictors, comparing five regressors under GroupKFold cross-validation keyed on the synthetic *farm + point* group (Fig. 5); Linear Regression attained the highest cross-validated skill (CV $R^2 = 0.95$), with Random Forest a distant second (CV $R^2 = 0.75$) and the remaining estimators well below. The $n = 71$ reflects a second, milder form of the incomplete coverage already visible in the audit trail: ten of the 81 points were surveyed with the chlorophyll meter and ceptometer only, with no IRGA or fluorometer readings, so A itself is undefined there and the fit uses the remaining 71 records (Farm A 27, Farm B 44); unlike the four points the empty-grid filter removes for carrying no measurement at all, these ten pass the pipeline legitimately and drop out only when a model of A is requested. This run illustrates the workflow rather than a physiological model: the two indices that contain A algebraically, WUE ($= A/E$) and A/C_i , were kept out of the predictor set, but several of the remaining predictors (g_s , C_i , C_i/C_a) are still mechanistically tied to A : under the gas-exchange equations used to obtain it, C_i itself takes A as an input, so the high skill largely restates that dependence and is not evidence of predictive performance on independent campaigns. The VIF caveat in the exploratory tab addresses a related but distinct problem, the mutual collinearity among the predictors; it does not measure their dependence on the response. With only three distinct collection dates (2025-12-19 to 2026-02-28), the STL guard-rail *blocked* the decomposition and explained why (Fig. 6), rather than returning a decomposition dominated by interpolated values.




Figure 5: Modelling page comparing five supervised regressors for net photosynthesis (*A*) on the Rio Verde dataset in a single run, reporting holdout R^2 , MAE and RMSE alongside cross-validated R^2 (mean and standard deviation). Linear Regression attains the highest cross-validated skill, with Random Forest a distant second; here the GroupKFold-per-site guard-rail (grouping on *farm + point*) is enabled to curb pseudoreplication.

STL decomposition (trend + seasonal + residual)

Robust decomposition of a univariate time series (Cleveland et al., 1990).

Variable

EUA

Seasonal period (days)

7

☒ Interpolate temporal gaps

Only 3 dates have actual measurements (minimum 10). STL decomposition needs enough observed points — sparse field campaigns do not meet the requirement.

Figure 6: Time Series page on the Rio Verde dataset. Because the campaign contains only three distinct collection dates, below the ten-date threshold, the STL guard-rail blocks the decomposition and reports the reason instead of producing an interpolation-dominated result.

3.2 Datasets beyond ecophysiology

To show that the platform is not bound to its original domain, we ran two public datasets through the same interface; in both, the profile resolver switched to the generic profile, so neither analysis required any change to the application.

The first is the classical Yates oats trial [42], three oat varieties under four nitrogen rates across six blocks (72 rows), uploaded as a plain CSV and analysed on the Experimental-design page as a split-plot, with variety as the whole-plot factor and nitrogen as the sub-plot factor. The resulting ANOVA agrees with the reference split-plot analysis of the same trial in R [43] (Table 2): the three F statistics agree to three decimal places. Variety is tested against the whole-plot error, whereas nitrogen and the variety $\times$ nitrogen interaction are tested against the sub-plot error; the variety effect is not statistically significant, nitrogen has a strong effect, and the interaction is not statistically significant. The post-hoc Tukey test on nitrogen recovers the textbook dose response and did not detect a statistically significant difference in yield between the two highest rates, and the page’s residual checks (Shapiro–Wilk, Levene) raise no flag.

Term (error stratum)	df	PhysioFlow F (p)	R reference F (p)
Variety (whole-plot)	2, 10	1.485 (0.272)	1.485 (0.272)
Nitrogen (sub-plot)	3, 45	37.686 (< 0.001)	37.686 (< 0.001)
Variety $\times$ Nitrogen	6, 45	0.303 (0.932)	0.303 (0.932)

Table 2: Split-plot ANOVA of the Yates oats trial computed by PhysioFlow and by a reference analysis in R. The F statistics agree to three decimal places for all three tested effects. Variety is tested against the whole-plot error (10 residual degrees of freedom); nitrogen and the interaction against the sub-plot error (45 residual degrees of freedom).

The second is the Palmer penguins survey [44], a biology dataset of 344 morphometric records from three species. The EDA recovered the known flipper-length/body-mass correlation ($r = 0.87$) and the sign reversal of bill depth against mass, a textbook instance of Simpson’s paradox. In classification mode, using the four morphometric predictors (bill length, bill depth, flipper length and body mass) and discarding records with a missing predictor or species label, Random Forest, SVM and k -NN each reached 1.00 accuracy on a stratified 25% holdout (`random_state = 42`) and 0.98–0.99 mean accuracy under five-fold stratified cross-validation, the latter computed on all eligible records rather than on the holdout split alone. These runs demonstrate the workflow on a public biological dataset; they are not a benchmark comparison under a common protocol. A split-plot agronomy trial and a biology survey thus run through the same guard-railed interface that served the gas-exchange campaign.

4 Impact overview

PhysioFlow brings the cleaning, exploratory analysis, supervised modelling and experimental-design analysis needed to interpret crop-ecophysiology data into a single platform. For researchers without a programming background, especially in under-resourced institutions, that barrier often delays the analysis of campaigns already expensive to run.

PhysioFlow’s main contribution is methodological, not just operational: its statistical guard-rails turn assumptions that researchers usually check informally, or skip, into explicit on-screen warnings, steering non-statistician users away from confounding, multicollinearity, pseudoreplication and sparsity. In the Rio Verde campaign they caught a perfectly confounded design and a series too short for seasonal decomposition at once, problems that, left unflagged, would more likely have surfaced only at peer review.

PhysioFlow is in routine use by two communities. Researchers of the Goiás Verde initiative (Instituto Federal Goiano and CEAGRE) apply it to soybean and sugarcane photosynthetic performance across Cerrado cropping systems, and graduate students of the Postgraduate Programme in Agricultural Sciences at Instituto Federal

Goiano, Campus Rio Verde, use it to analyse their own experiments and as a teaching tool in data-analysis training. Because these users span plant physiology, agronomy and soil science, the same guard-railed interface serves both the research and the training pipeline of the group. The current release (v1.0) comprises about 7,500 lines of Python, 169 `pytest` test cases and 775 user-interface keys. Informal user feedback reports shorter analysis times than the spreadsheet-based workflow it replaced; no controlled evaluation of time saving has been carried out. No publication has yet been produced with *PhysioFlow*: the present evidence of use comprises the research and teaching activities described above, and journal manuscripts based on ongoing analyses are in preparation (Section 6).

5 Limitations and future development

PhysioFlow has known limitations. It targets tabular exports and does not read proprietary binary formats, so vendor files must first be exported to CSV, TSV or XLSX. The physiology profile recognises a fixed 31-column reference schema; datasets outside it fall back to the generic profile, where domain presets and instrument-specific replicate handling are unavailable. The experimental-design module assumes balanced or nearly balanced layouts, and it reports, but does not automatically correct, violated assumptions; severely unbalanced or missing-cell designs still require a dedicated mixed-model package. The current Dunnett implementation is restricted to one-way comparisons and does not incorporate blocks or design-specific error strata, and the external comparison against `multcomp` was likewise confined to one-way examples; its results should therefore not be read as design-adjusted comparisons in blocked, factorial or split-plot experiments. Finally, the application is session-based and holds the working dataset in memory, which bounds the size of campaigns it can process interactively on a standard laptop.

Several guard-rails apply beyond ecophysiology when their data requirements are met, but they are of different kinds and should not be read as a single guarantee. The categorical association measure (bias-corrected Cramér’s V [28]), the multicollinearity diagnostic (VIF) and the residual normality and homoscedasticity checks (Shapiro–Wilk, Levene) address distinct aspects of statistical diagnosis, whereas GroupKFold is a partitioning strategy that keeps user-defined groups separate across cross-validation folds. The ten-date threshold for STL is an application-level safeguard against sparse series, not a universal criterion for a valid seasonal decomposition, which also depends on the sampling interval, the seasonal period, missingness and the number of observed cycles. Defaults and interpretation remain domain-dependent: the physiology profile uses gas-exchange presets and a synthetic *farm+point* grouping key and assumes the 31-column schema, whereas users of the generic profile must themselves identify the relevant factors, grouping units and sampling structure. The Yates oats and Palmer penguins examples demonstrate use beyond ecophysiology, but do not establish validity for every tabular dataset. These safeguards flag specific risks; they do not replace domain expertise, nor do they guarantee detection of measurement error, instrument drift or every form of dependence.

Planned developments include physiological range validation in the schema (`valid_range` per column), integration of the instrument data dictionary as in-app tooltips, raising test coverage above 70% across the `ml/`, `state` and `auth` modules, a quasi-constant-variable warning in the correlation panel, and a public Docker image for one-command deployment. On the analytical side, automatic fitting of physiological response curves (e.g. A/C_i and light-response curves), automated report export, and predictive guard-rails that flag physiological-stress patterns to support decision-making are planned as future directions.

6 Publications enabled by the software

No publication has yet been produced with *PhysioFlow*. The platform was presented at the 1st International Symposium on Climate and Agricultural Resilience (5th CEAGRE Agro Experience, Rio Verde, 2026), and it is in routine use in the ongoing analyses of the Goiás Verde initiative and in the graduate work of the Postgraduate Programme in Agricultural Sciences at Instituto Federal Goiano, Campus Rio Verde; studies on soybean and sugarcane photosynthetic performance in the Cerrado are in preparation, and this section will grow as those

campaigns are published. The abstract published in the proceedings of that meeting [19] describes an earlier platform by the same group, for greenhouse-gas flux data, which motivated the present work (Section 1); it was not produced with *PhysioFlow*.

CRedit authorship contribution statement

Daniel Hilário da Silva: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Visualization, Writing – original draft, Writing – review & editing. **Leandro Rodrigues da Silva Souza:** Conceptualization, Methodology, Software, Supervision, Project administration, Funding acquisition, Writing – review & editing. **André Abade:** Conceptualization, Methodology, Software, Validation, Writing – review & editing. **Francisco de Almeida Lobo:** Validation, Supervision, Writing – review & editing. **Fernando Rodrigues Cabral Filho:** Investigation, Data curation, Writing – review & editing. **Erli Pinto dos Santos:** Software, Visualization, Writing – review & editing. **Adinan Alves da Silva:** Investigation, Data curation, Resources. **Daiane Alves da Silva:** Investigation, Resources, Writing – review & editing. **Ketlyn Santos Sousa:** Investigation, Data curation, Writing – review & editing. **Igor Eli da Silva:** Software, Validation, Visualization. **Alan Carlos da Costa:** Supervision, Funding acquisition, Resources, Writing – review & editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The source code and the documentation are available at the version-specific repository link given in C2, and the public datasets used in Section 3.2 (Yates oats and Palmer penguins) ship under `data/sample/test/`. After installing the dependencies, running `python reproduce_examples.py` from the repository root regenerates the split-plot ANOVA of Table 2 and the penguins correlation and classification results. The published Code Ocean capsule (C3) archives the computational environment, the scripts and the outputs of its reproducible run, which covers the public examples, the comparisons against R and the `pytest` suite; a Streamlit workstation opens the application itself. The capsule pins commit 9960397; the repository link in C2 points to a later commit, which adds the post-hoc comparisons against R and their documentation, plus lint, type-checking and continuous-integration fixes that leave the numerical results unchanged. The field measurements of Section 3.1 belong to an ongoing research project and cannot be released; the figures derived from them are de-identified. A synthetic ecophysiology dataset reproduces the reference schema and demonstrates the cleaning workflow and the guard-rail behaviour discussed in Section 3.1, including the categorical confounding and the sparse temporal sampling, but it does not reproduce the numerical results of that campaign.

Acknowledgements

The authors thank the Ministry of Science, Technology and Innovation (MCTI), the Brazilian Funding Authority for Studies and Projects (FINEP), the Goiás State Research Foundation (FAPEG), the National Council for Scientific and Technological Development (CNPq), the Coordination for the Improvement of Higher Education Personnel (CAPES), the State Secretariat for Development and Innovation of Goiás (SECTI), the Center of Excellence in Exponential Agriculture (CEAGRE), and the Instituto Federal Goiano for supporting the research developed.

Declaration of generative AI in scientific writing

During the preparation of this work the authors used an AI assistant (Claude, Anthropic) to verify the consistency between the values reported in the text and those shown in the figures, to assist with debugging the software, and to improve the wording of specific passages. The figures are unaltered screen captures of the application in use, produced by a script that drives the interface and records the resulting screens; the de-identification described in Section 3.1 was applied to the input data, not to the images. The authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

References

- [1] Hans Lambers, F. Stuart Chapin, and Thijs L. Pons. *Plant Physiological Ecology*. Springer, New York, 2nd edition, 2008. doi: 10.1007/978-0-387-78341-3.
- [2] Stephen P. Long, Xin-Guang Zhu, Shawna L. Naidu, and Donald R. Ort. Can improvement in photosynthesis increase crop yields? *Plant, Cell & Environment*, 29(3):315–330, 2006. doi: 10.1111/j.1365-3040.2005.01493.x.
- [3] G. D. Farquhar, S. von Caemmerer, and J. A. Berry. A biochemical model of photosynthetic CO_2 assimilation in leaves of C_3 species. *Planta*, 149(1):78–90, 1980. doi: 10.1007/BF00386231.
- [4] Susanne von Caemmerer. *Biochemical Models of Leaf Photosynthesis*. CSIRO Publishing, Collingwood, 2000.
- [5] S. P. Long and C. J. Bernacchi. Gas exchange measurements, what can they tell us about the underlying limitations to photosynthesis? procedures and sources of error. *Journal of Experimental Botany*, 54(392):2393–2401, 2003. doi: 10.1093/jxb/erg262.
- [6] Kate Maxwell and Giles N. Johnson. Chlorophyll fluorescence—a practical guide. *Journal of Experimental Botany*, 51(345):659–668, 2000. doi: 10.1093/jexbot/51.345.659.
- [7] Neil R. Baker. Chlorophyll fluorescence: a probe of photosynthesis in vivo. *Annual Review of Plant Biology*, 59:89–113, 2008. doi: 10.1146/annurev.arplant.59.032607.092759.
- [8] Nathalie J. J. Bréda. Ground-based measurements of leaf area index: a review of methods, instruments and current controversies. *Journal of Experimental Botany*, 54(392):2403–2417, 2003. doi: 10.1093/jxb/erg263.
- [9] Wanderson J. de Oliveira, Edivan R. de Souza, Jailson C. Cunha, Ênio F. de F. e Silva, and Venâncio de L. Veloso. Leaf gas exchange in cowpea and CO_2 efflux in soil irrigated with saline water. *Revista Brasileira de Engenharia Agrícola e Ambiental*, 21(1):32–37, 2017. doi: 10.1590/1807-1929/agriambi.v21n1p32-37.
- [10] L. Munjonji, K. K. Ayisi, T. P. Mafeo, T. Maphanga, and K. E. Mabitsele. Seasonal variation in soil CO_2 emission and leaf gas exchange of well managed commercial citrus sinensis (l.) orchards. *Plant and Soil*, 465:65–81, 2021. doi: 10.1007/s11104-021-04986-x.
- [11] Emily B. Peters and Joseph P. McFadden. Continuous measurements of net CO_2 exchange by vegetation and soils in a suburban landscape. *Journal of Geophysical Research: Biogeosciences*, 117:G03005, 2012. doi: 10.1029/2011JG001933.
- [12] M. L. Sánchez, M. I. Ozores, M. J. López, R. Colle, B. De Torre, M. A. García, and I. Pérez. Soil CO_2 fluxes beneath barley on the central spanish plateau. *Agricultural and Forest Meteorology*, 118:85–95, 2003. doi: 10.1016/S0168-1923(03)00098-3.
- [13] S. Guarrera, D. Vanella, S. Consoli, G. Giudice, S. Toscano, J. M. Ramírez-Cuesta, M. Milani, F. Ferlito, and D. Longo. Analysis of small-scale soil CO_2 fluxes in an orange orchard under irrigation and soil conservative practices. *Heliyon*, 10(9):e30543, 2024. doi: 10.1016/j.heliyon.2024.e30543.

- [14] George L. Vourlitis, Osvaldo Borges Pinto Jr., Higo J. Dalmagro, Paulo Enrique Zanella de Arruda, Francisco de Almeida Lobo, and José de Souza Nogueira. Net primary production and ecosystem carbon flux of brazilian tropical savanna ecosystems from eddy covariance and inventory methods. *Journal of Geophysical Research: Biogeosciences*, 127(3):e2021JG006780, 2022. doi: 10.1029/2021JG006780.
- [15] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2024. URL <https://www.R-project.org/>.
- [16] Stuart H. Hurlbert. Pseudoreplication and the design of ecological field experiments. *Ecological Monographs*, 54(2):187–211, 1984. doi: 10.2307/1942661.
- [17] Streamlit Inc. Streamlit: A faster way to build and share data apps. <https://streamlit.io>, 2024.
- [18] Daniel Hilário da Silva, Leandro Rodrigues da Silva Souza, Caio Tonus Ribeiro, Simone Hilário da Silva Brasileiro, José Renato Munari Nardo, Adriano Alves Pereira, and Adriano de Oliveira Andrade. A web application for exploratory data analysis and classification of Parkinson’s disease patients using machine learning models on different datasets. *Software Impacts*, 23:100737, 2025. doi: 10.1016/j.simpa.2024.100737.
- [19] Daniel Hilário da Silva, André Abade, André Thomazini, Leandro Rodrigues da Silva Souza, Fernando Rodrigues Cabral Filho, Tiago do Prado Paim, Erli Pinto dos Santos, and Douglas Farias Cordeiro. Interactive web platform for processing and predictive modeling of greenhouse gas fluxes in soil. In Alan Carlos da Costa, Fernando Rodrigues Cabral Filho, Ícaro Lunas Nunes, and Leandro Rodrigues da Silva Souza, editors, *Anais de Resumos do 5º CEAGRE Agro Experience, 4º Dia de Campo do LEAV e 1st International Symposium on Climate and Agricultural Resilience (ISCAR)*, Rio Verde, Goiás, Brazil, 2026. Instituto Federal Goiano. ISBN 978-85-65871-77-8. Abstract.
- [20] Richard Magala and Lisa A. Schulte. apsimngpy: A comprehensive python framework for interactive, reproducible, and scalable simulations of the apsim next generation model. *SoftwareX*, 33:102496, 2026. doi: 10.1016/j.softx.2025.102496.
- [21] Lorenzo Crecco and Sofia Bajocco. prismatools: An open-source python package for accessing and analyzing prisma hyperspectral data. *SoftwareX*, 34:102547, 2026. doi: 10.1016/j.softx.2026.102547.
- [22] Wes McKinney. Data structures for statistical computing in python. In *Proceedings of the 9th Python in Science Conference*, pages 56–61, 2010. doi: 10.25080/Majora-92bf1922-00a.
- [23] Charles R. Harris et al. Array programming with numpy. *Nature*, 585:357–362, 2020. doi: 10.1038/s41586-020-2649-2.
- [24] F. Pedregosa et al. Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12: 2825–2830, 2011.
- [25] Michael L. Waskom. seaborn: statistical data visualization. *Journal of Open Source Software*, 6(60):3021, 2021. doi: 10.21105/joss.03021.
- [26] John D. Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3): 90–95, 2007. doi: 10.1109/MCSE.2007.55.
- [27] Sergio J. Rey and Luc Anselin. Pysal: A python library of spatial analytical methods. *The Review of Regional Studies*, 37(1):5–27, 2007.
- [28] Wicher Bergsma. A bias-correction for cramér’s v and tschuprow’s t. *Journal of the Korean Statistical Society*, 42(3):323–328, 2013. doi: 10.1016/j.jkss.2012.10.002.
- [29] Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001. doi: 10.1023/A:1010933404324.
- [30] Jerome H. Friedman. Greedy function approximation: a gradient boosting machine. *The Annals of Statistics*, 29(5):1189–1232, 2001. doi: 10.1214/aos/1013203451.

- [31] William H. Kruskal and W. Allen Wallis. Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260):583–621, 1952. doi: 10.1080/01621459.1952.10483441.
- [32] S. S. Shapiro and M. B. Wilk. An analysis of variance test for normality (complete samples). *Biometrika*, 52(3–4):591–611, 1965. doi: 10.1093/biomet/52.3-4.591.
- [33] David A. Belsley, Edwin Kuh, and Roy E. Welsch. *Regression Diagnostics: Identifying Influential Data and Sources of Collinearity*. John Wiley & Sons, New York, 1980. doi: 10.1002/0471725153.
- [34] Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. Isolation forest. In *2008 Eighth IEEE International Conference on Data Mining*, pages 413–422, 2008. doi: 10.1109/ICDM.2008.17.
- [35] Markus M. Breunig, Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. Lof: Identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*, pages 93–104, 2000. doi: 10.1145/342009.335388.
- [36] A. J. Scott and M. Knott. A cluster analysis method for grouping means in the analysis of variance. *Biometrics*, 30(3):507–512, 1974. doi: 10.2307/2529204.
- [37] Luc Anselin. Local indicators of spatial association—LISA. *Geographical Analysis*, 27(2):93–115, 1995. doi: 10.1111/j.1538-4632.1995.tb00338.x.
- [38] Arthur Getis and J. K. Ord. The analysis of spatial association by use of distance statistics. *Geographical Analysis*, 24(3):189–206, 1992. doi: 10.1111/j.1538-4632.1992.tb00261.x.
- [39] Noel A. C. Cressie. *Statistics for Spatial Data*. John Wiley & Sons, New York, revised edition, 1993. doi: 10.1002/9781119115151.
- [40] Robert B. Cleveland, William S. Cleveland, Jean E. McRae, and Irma Terpenning. Stl: A seasonal-trend decomposition procedure based on loess. *Journal of Official Statistics*, 6(1):3–73, 1990.
- [41] H. B. Mann and D. R. Whitney. On a test of whether one of two random variables is stochastically larger than the other. *The Annals of Mathematical Statistics*, 18(1):50–60, 1947. doi: 10.1214/aoms/1177730491.
- [42] Frank Yates. Complex experiments. *Supplement to the Journal of the Royal Statistical Society*, 2(2): 181–247, 1935. doi: 10.2307/2983638.
- [43] José C. Pinheiro and Douglas M. Bates. *Mixed-Effects Models in S and S-PLUS*. Springer, New York, 2000. doi: 10.1007/978-1-4419-0318-1.
- [44] Kristen B. Gorman, Tony D. Williams, and William R. Fraser. Ecological sexual dimorphism and environmental variability within a community of antarctic penguins (genus *Pygoscelis*). *PLoS ONE*, 9(3):e90081, 2014. doi: 10.1371/journal.pone.0090081.